

Daiwei Chen

daiwei.chen@wisc.edu | +1 445-208-7806 | Website | LinkedIn | Google Scholar

RESEARCH FOCUS: LLM FACTUALITY, SAFETY & EVALUATION

Ph.D. researcher studying LLM factuality, robustness, and the safety-capability trade-off. Research spans RAG evaluation, pluralistic alignment, and LinkedIn Responsible AI, with an emphasis on both reliability and usefulness.

EDUCATION

University of Wisconsin-Madison

Ph.D. in Electrical and Computer Engineering

Advisor: Ramya Korlakai Vinayak

Madison, WI

Sep 2023 – Present

University of Pennsylvania

M.S. in Systems Engineering

Philadelphia, PA

Aug 2021 – May 2023

Zhejiang University

B.S. in Psychology

Hangzhou, China

Sep 2017 – Jun 2021

INDUSTRY RESEARCH EXPERIENCE

LinkedIn

Mountain View, CA

Responsible AI - Research Intern

Jun 2026 – Sep 2026

- Developed **Multi-Teacher On-Policy Distillation (MOPD)** to address the **alignment tax**: the trade-off between safety alignment and model capabilities in LLM post-training.

Generative AI - Research Intern

May 2025 – Aug 2025

- Developed **GTI** for generative candidate retrieval; improved offline **Precision@5 by 21.6% relative** to vanilla SFT under the Good Match criterion on an industrial dataset.

SELECTED RESEARCH & PUBLICATIONS

Is Conformal Factuality for RAG-based LLMs Robust? Novel Metrics and Systematic Insights

Y. Chen, [D. Chen](#), S. M. Chikodikar, C. H. Yin, R. K. Vinayak

Preprint, arXiv:2603.16817, 2026

Utility-aware metrics reveal when factuality filtering sacrifices useful information or fails under shift.

One Model, Many Modes: Evolving System Prompt Ensembles for Creative Generation

[D. Chen](#), A. He, B. Yao, J. Hu, R. K. Vinayak

Under review

Prompt ensembles improve novelty and diversity while preserving quality across four tasks and three models.

PAL: Sample-Efficient Personalized Reward Modeling for Pluralistic Alignment

D. Chen, Y. Chen, A. Rege, Z. Wang, R. K. Vinayak

ICLR 2025 | NeurIPS 2024 Workshop (Spotlight) | ICML 2024 Workshop (Oral)

Personalized reward models represent diverse human preferences and adapt to new users with few examples.

Grounded Token Initialization for New Vocabulary in LMs for Generative Recommendation

[D. Chen](#), Z. Fu, C. Jiang, H. Zhang, R. Zhou, T. Wang, C. Yao, G. Li, R. Cai, Y. Cao, R. Jiang, F. Borisyuk, J. Shen, J. Wu, R. K. Vinayak

EMNLP 2026 Main Conference

Spectral diagnostics reveal token-embedding collapse; grounded initialization improves generative retrieval.

Compositional Self-Improvement

C. Shin, [D. Chen](#), J. Cooper, B. Lake, F. Sala, R. K. Vinayak

ICML 2026 – 2nd Workshop on Compositional Learning: Safety, Interpretability, and Agents

Composing reliable subproblem predictions enables self-training beyond the original supervision regime.

SERVICE & AWARDS

Reviewer: NeurIPS 2026, ICLR 2026, CVPR 2026, EACL 2026, COLM 2026.

Awards: Lambda.ai Research Grant Sponsorship (2025); UPenn Outstanding Research Award (2023).