

Daiwei Chen

daiwei.chen@wisc.edu | +1 445-208-7806 | Website | LinkedIn | Google Scholar

RESEARCH FOCUS: LLM POST-TRAINING, PREFERENCE LEARNING & ALIGNMENT

Ph.D. researcher working on personalized reward modeling, pluralistic alignment, and the safety-capability trade-off in LLM post-training. First-author PAL (ICLR 2025); LinkedIn Responsible AI and GenAI research experience.

EDUCATION

University of Wisconsin-Madison

Ph.D. in Electrical and Computer Engineering

Advisor: Ramya Korlakai Vinayak

Madison, WI

Sep 2023 – Present

University of Pennsylvania

M.S. in Systems Engineering

Philadelphia, PA

Aug 2021 – May 2023

Zhejiang University

B.S. in Psychology

Hangzhou, China

Sep 2017 – Jun 2021

INDUSTRY RESEARCH EXPERIENCE

LinkedIn

Mountain View, CA

Responsible AI - Research Intern

Jun 2026 – Sep 2026

- Developed **Multi-Teacher On-Policy Distillation (MOPD)** to address the **alignment tax** in LLM post-training, targeting safety alignment while preserving model capabilities.

Generative AI - Research Intern

May 2025 – Aug 2025

- Developed **GTI** for generative candidate retrieval; improved offline **Precision@5 by 21.6% relative** to vanilla SFT under the Good Match criterion on an industrial dataset.

SELECTED RESEARCH & PUBLICATIONS

PAL: Sample-Efficient Personalized Reward Modeling for Pluralistic Alignment

D. Chen, Y. Chen, A. Rege, Z. Wang, R. K. Vinayak

ICLR 2025 | NeurIPS 2024 Workshop (*Spotlight*) | ICML 2024 Workshop (*Oral*)

Shared preference prototypes enable sample-efficient reward modeling and few-shot adaptation to new users.

One Model, Many Modes: Evolving System Prompt Ensembles for Creative Generation

D. Chen, A. He, B. Yao, J. Hu, R. K. Vinayak

Under review

Evolved system-prompt ensembles improve novelty and diversity without retraining or sacrificing quality.

Compositional Self-Improvement

C. Shin, D. Chen, J. Cooper, B. Lake, F. Sala, R. K. Vinayak

ICML 2026 – 2nd Workshop on Compositional Learning: Safety, Interpretability, and Agents

Composing reliable subproblem predictions enables self-training beyond the original supervision regime.

Is Conformal Factuality for RAG-based LLMs Robust? Novel Metrics and Systematic Insights

Y. Chen, D. Chen, S. M. Chikodikar, C. H. Yin, R. K. Vinayak

Preprint, arXiv:2603.16817, 2026

Utility-aware evaluation exposes factuality filtering failures under distribution shift and distractors.

Grounded Token Initialization for New Vocabulary in LMs for Generative Recommendation

D. Chen, Z. Fu, C. Jiang, H. Zhang, R. Zhou, T. Wang, C. Yao, G. Li, R. Cai, Y. Cao, R. Jiang, F. Borisyuk, J. Shen, J. Wu, R. K. Vinayak

EMNLP 2026 Main Conference

Grounding new token embeddings before fine-tuning improves generative retrieval and recommendation.

SERVICE & AWARDS

Reviewer: NeurIPS 2026, ICLR 2026, CVPR 2026, EACL 2026, COLM 2026.

Awards: Lambda.ai Research Grant Sponsorship (2025); UPenn Outstanding Research Award (2023).