Daiwei Chen

daiwei.chen@wisc.edu | +1 445-208-7806 | Website | LinkedIn | Google Scholar

Research Focus: LLM Post-Training, Preference Learning & Alignment

Ph.D. researcher working on personalized reward modeling, pluralistic alignment, and the safety–capability trade-off in LLM post-training. First-author PAL (**ICLR 2025**); LinkedIn Responsible AI and GenAI research experience.

Education

University of Wisconsin-Madison — Madison, WI
Ph.D. in Electrical and Computer Engineering — *Sep 2023 – Present*
Advisor: Ramya Korlakai Vinayak

University of Pennsylvania — Philadelphia, PA
M.S. in Systems Engineering — *Aug 2021 – May 2023*

Zhejiang University — Hangzhou, China
B.S. in Psychology — *Sep 2017 – Jun 2021*

Industry Research Experience

LinkedIn — Mountain View, CA

Responsible AI - Research Intern — Jun 2026 – Sep 2026

- Developed **Multi-Teacher On-Policy Distillation (MOPD)** to address the **alignment tax** in LLM post-training, targeting safety alignment while preserving model capabilities.

Generative AI - Research Intern — May 2025 – Aug 2025

- Developed **GTI** for generative candidate retrieval; improved offline **Precision@5 by 21.6% relative** to vanilla SFT under the Good Match criterion on an industrial dataset.

Selected Research & Publications

PAL: Sample-Efficient Personalized Reward Modeling for Pluralistic Alignment
D. Chen, Y. Chen, A. Rege, Z. Wang, R. K. Vinayak
ICLR 2025 | *NeurIPS 2024 Workshop (**Spotlight**)* | *ICML 2024 Workshop (**Oral**)*
Shared preference prototypes enable sample-efficient reward modeling and few-shot adaptation to new users.

One Model, Many Modes: Evolving System Prompt Ensembles for Creative Generation
D. Chen, A. He, B. Yao, J. Hu, R. K. Vinayak
Under review
Evolved system-prompt ensembles improve novelty and diversity without retraining or sacrificing quality.

Compositional Self-Improvement
C. Shin, D. Chen, J. Cooper, B. Lake, F. Sala, R. K. Vinayak
ICML 2026 *– 2nd Workshop on Compositional Learning: Safety, Interpretability, and Agents*
Composing reliable subproblem predictions enables self-training beyond the original supervision regime.

Is Conformal Factuality for RAG-based LLMs Robust? Novel Metrics and Systematic Insights
Y. Chen, D. Chen, S. M. Chikodikar, C. H. Yin, R. K. Vinayak
Under review | *Preprint, arXiv:2603.16817, 2026*
Utility-aware evaluation exposes factuality filtering failures under distribution shift and distractors.

Grounded Token Initialization for New Vocabulary in LMs for Generative Recommendation
D. Chen, Z. Fu, C. Jiang, H. Zhang, R. Zhou, T. Wang, C. Yao, G. Li, R. Cai, Y. Cao, R. Jiang, F. Borisyuk, J. Shen, J. Wu, R. K. Vinayak
EMNLP 2026 *Main Conference*
Grounding new token embeddings before fine-tuning improves generative retrieval and recommendation.

Service & Awards

Reviewer: NeurIPS 2026, ICLR 2026, CVPR 2026, EACL 2026, COLM 2026.
Awards: Lambda.ai Research Grant Sponsorship (2025); UPenn Outstanding Research Award (2023).